

Progress Report

Who worked on each part of the design

Jiamao Xu: Local branch history table, Global 2-level branch history table, tournament branch prediction

Jerry Wang: L2+ cache system, basic and strided prefetching

Garen Hu: Parameterized cache

The functionalities you implemented

Our cache module has undergone a significant upgrade. We have implemented a two-level cache structure, which includes a level 1 (L1) cache and a level 2 (L2) cache. This allows us to achieve shorter memory response times, as the processor can access the L1 cache faster than the main memory. Moreover, the L2 cache acts as a buffer between the main memory and the L1 cache, thereby further reducing access time. Additionally, we have incorporated up to 8-way set-associative cache systems to improve the overall caching performance. This means that our cache system can store data in multiple ways, which results in fewer cache misses and improved system performance.

To ensure our cache system is more flexible, we have added a parameterized cache. This allows us to easily change the cache size and associativity, depending on the requirements of the workload. This feature enhances the overall performance of the CPU, as it adapts to varying workloads and data access patterns.

Another advanced feature we have incorporated is a branch prediction module. Branches are an integral part of program execution, but they can also result in a significant time penalty. To minimize this, we have implemented two different strategies to improve our prediction accuracy. The first strategy is local branch history table (LBHT) prediction, which uses the branch history of the current program execution to predict the outcome of a branch. The second strategy is global two-level branch history table (BHT) prediction, which stores the history of all branches in the program and uses this information to predict the outcome of a branch.

Overall, these advanced features improve the performance of our CPU by reducing the time penalty associated with cache misses and branch operations, while also ensuring flexibility in handling different workloads.

The testing strategy you used to verify these functionalities

To ensure that our CPU is functioning correctly and passing the test code for cp3, we executed the provided test code and used Verdi to verify the memory address and the corresponding register value. If we encountered any issues or inconsistencies, we manually located the memory address causing the problem and adjusted the halting

condition to a few lines after the problematic line. This allowed us to troubleshoot any issues and ensure the basic functionality of the CPU is operating correctly.

For the advanced features, we conducted additional testing to ensure their effectiveness. We checked the registers to see if they contained the correct value when we changed the cache settings. This allowed us to verify that the cache was functioning correctly and storing and retrieving data as expected. We also compared the time costs between running programs with and without advanced cache features to determine the extent to which the cache improved performance.

Furthermore, we incorporated a counter register to evaluate the accuracy of our branch prediction system. By keeping track of the number of correct and incorrect predictions, we were able to assess the overall effectiveness of our branch prediction strategies. This helped us identify areas for improvement and refine our branch prediction system to maximize performance.

Overall, our testing and verification process involved multiple steps and techniques to ensure the basic and advanced features of our CPU were functioning correctly. By thoroughly testing and evaluating each feature, we were able to identify and troubleshoot any issues and optimize the performance of our CPU.

Road Map

Who is going to implement and verify each feature or functionality you must complete

Jiamao Xu: Slides, design competition basic and competition coremark

Zihan Hu: Slides, design competition basic and competition coremark

Jerry Wang: Slides, design competition basic and competition coremark

What are those features or functionalities

Our first priority is to finish designing our slides for the presentation, which will include detailed information on our 5-stage pipeline CPU and the advanced features we have incorporated. Each team member will be responsible for designing a couple of pages and preparing them for the presentation. This will ensure that we have a comprehensive and well-organized presentation that covers all aspects of our CPU design and implementation.

After completing the slide design, we will spend time practicing the presentation as a team. We will run through the presentation multiple times, making adjustments as necessary to ensure that our delivery is smooth, concise, and engaging. Additionally, we

will work on improving our public speaking skills and addressing any potential questions or concerns that may arise during the presentation.

For the design competition is to ensure that we complete the competition's basic requirements and obtain the necessary points. To achieve this, we will focus on building a functional and reliable CPU that meets all the specified requirements. We will work closely together as a team to ensure that each member is contributing their skills and expertise to the design process.

Once we have completed the basic requirements of the competition, we will assess our progress and determine if we have any additional time to improve our design performance. If we do, we will focus on optimizing our CPU's performance, improving its efficiency, and enhancing its features.

Progress Report

Who worked on each part of the design

Jiamao Xu: Hazard detection & forwarding & Static branch predictor & RVFI

Garen Hu: Arbiter Hazard detection & forwarding & Static branch predictor

Jerry Wang: Integration of L1 cache Hazard detection & forwarding & RVFI

The functionalities you implemented

Based on the function we came up with from CP1, we added more advanced implementations including L1 cache system, hazard detection & forwarding, and static branch predictor. We also integrate the RVFI monitor to compare with the true result.

The testing strategy you used to verify these functionalities

First, we run the provided test code, and we use Verdi to verify the mem address and the corresponding register value. If we find something goes wrong. We find the mem address and we manually change the halting condition to a few lines after the problem line. For example, if the problem occurs at memory address 80000198, we set up halting at 800001a0 to prevent executing code after it. And then we start with the instruction data and instruction address and trace the data along the datapath to see where the problem is. In addition, we spent a huge time debugging our forwarding unit, we kept tracing the key registers from the end to where it went wrong and tried to fix our code.

The timing and energy analysis of your design: fmax & energy report from Design Compiler

```

clock my_clk (rise edge)          10.00    10.00
clock network delay (ideal)        0.00    10.00
clock uncertainty                   -0.10     9.90
cpu/datapath/PC_register/data_reg[31]/CK (DFF_X1) 0.00     9.90 r
library setup time                  -0.04     9.86
data required time                  9.86
-----
data required time                  9.86
data arrival time                   -5.78
-----
slack (MET)                        4.08

```

$f_{\max} = 1 / (10 - \text{slack}) * 1000 = 1 / (10 - 4.08) * 1000 = 168.918 \text{ Mhz}$

Power report:

Hierarchy	Switch Power	Int Power	Leak Power	Total Power	%
mp4	190.017	554.036	7.68e+05	1.51e+03	100.0
cache_adaptor (cacheline_adaptor)	17.178	33.625	3.66e+04	87.376	5.8
arbiter (arbiter)	5.157	2.041	2.34e+04	30.609	2.0
d_cache (cache_0)	69.017	209.797	2.44e+05	522.419	34.5
bus (line_adapter_0)	0.813	0.246	4.95e+03	6.009	0.4
datapath (cache_datapath_0)	68.061	209.092	2.38e+05	515.462	34.1
dirty (array_width1_0)	0.601	2.562	1.57e+03	4.729	0.3
valid (array_width1_1)	0.197	0.765	1.50e+03	2.464	0.2
tag (array_width24_0)	0.766	5.713	1.93e+04	25.801	1.7
DM_cache (data_array_0)	65.032	199.081	1.96e+05	460.384	30.4
control (cache_control_0)	0.143	0.459	345.840	0.948	0.1
i_cache (cache_1)	68.656	213.908	2.37e+05	519.712	34.4
bus (line_adapter_1)	1.408	0.901	4.60e+03	6.911	0.5
datapath (cache_datapath_1)	67.203	212.590	2.32e+05	512.190	33.9
valid (array_width1_3)	0.109	0.779	1.58e+03	2.466	0.2
tag (array_width24_1)	0.874	6.022	1.94e+04	26.267	1.7
DM_cache (data_array_1)	64.045	204.845	2.03e+05	471.973	31.2
control (cache_control_1)	4.39e-02	0.417	148.749	0.610	0.0
cpu (cpu)	29.331	94.459	2.26e+05	349.565	23.1
hzard_control (hazard_control)	4.38e-03	4.97e-03	45.430	5.48e-02	0.0
datapath (datapath)	28.844	94.171	2.23e+05	346.447	22.9
forward (forwarding)	0.352	0.378	1.75e+03	2.476	0.2
CMP (branch)	0.273	0.295	4.29e+03	4.855	0.3
ALU (alu)	5.165	4.221	2.26e+04	32.000	2.1
regfile (regfile)	1.954	22.658	1.11e+05	135.225	8.9
MEM_WB_mem_rdata (register_3)	0.208	1.083	3.24e+03	4.535	0.3
MEM_WB_aluout (register_4)	0.927	3.328	3.21e+03	7.466	0.5
MEM_WB_br_en (register_width1_0)	2.90e-02	0.445	147.346	0.621	0.0
MEM_WB_u_imm (register_5)	0.578	2.403	2.07e+03	5.046	0.3
MEM_WB_pc (pc_register_1)	0.946	3.374	3.27e+03	7.586	0.5
MEM_WB_ctrl (register_width46_0)	0.376	1.548	1.09e+03	3.019	0.2
EX_MEM_wdata (register_7)	0.867	3.317	3.16e+03	7.341	0.5
EX_MEM_br_en (register_width1_1)	2.99e-02	0.450	144.320	0.624	0.0
EX_MEM_u_imm (register_10)	0.563	2.460	2.07e+03	5.097	0.3
EX_MEM_aluout (register_11)	1.401	3.352	3.20e+03	7.953	0.5
EX_MEM_pc (pc_register_3)	0.892	3.284	3.23e+03	7.411	0.5
EX_MEM_ctrl (register_width46_1)	0.450	1.940	1.68e+03	4.069	0.3
ID_EX_rs2 (register_13)	0.977	3.682	3.19e+03	7.848	0.5
ID_EX_rsl (register_14)	0.971	2.965	3.25e+03	7.186	0.5
ID_EX_u_imm (register_15)	0.604	2.497	2.07e+03	5.169	0.3
ID_EX_i_imm (register_16)	0.913	3.684	2.76e+03	7.360	0.5
ID_EX_alumux2 (register_17)	1.006	3.736	3.19e+03	7.931	0.5
ID_EX_alumux1 (register_18)	0.998	3.107	3.26e+03	7.369	0.5
ID_EX_pc (pc_register_5)	0.972	3.554	3.24e+03	7.762	0.5
ID_EX_ctrl (register_width46_2)	1.347	4.190	3.99e+03	9.527	0.6
IF_ID_pc (pc_register_7)	0.990	3.556	3.23e+03	7.779	0.5
imm_decoder (ir)	1.368	3.807	3.26e+03	8.433	0.6
PC_register (pc_register_8)	1.160	3.317	3.23e+03	7.711	0.5
control (control_rom)	0.483	0.284	2.30e+03	3.064	0.2

Who is going to implement and verify each feature or functionality you must complete

Jiamao Xu: Local Branch History Table, Basic Hardware Prefetching

Zihan Hu: L2 cache, Parameterized cache

Jerry Wang: Victim cache, 4-way set associative cache

What are those features or functionalities

L2 cache system: L2 cache is a type of memory cache to improve our system performance. Since large L1 caches will form a large critical path, we need a L2 cache to prevent such issues

4-way set associative cache(may be 8-way set associative cache): we will implement a 4-way or 8-way set associative cache with pseudo-LRU replacement policy to make our cache system more efficient and have large cache size.

Parameterized cache: Parameterized cache is a cache memory design that allows for flexibility in its configuration by allowing users to specify the cache parameters such as cache size, block size, and associativity.

Victim Cache: A victim cache is a small, fully associative cache that stores blocks that have been evicted from a larger, more frequently accessed cache, known as the primary cache. When a cache miss occurs in the primary cache, the block that is evicted from the cache is stored in the victim cache.

Local Branch History Table:

The 2-bit local branch history table is a component of a branch predictor used in computer processors to improve performance. It stores information about the outcome of recent branch instructions and is used to predict whether a branch is likely to be taken or not taken in the future.

Basic Hardware Prefetching:

The OBL (One block lookahead) prefetch is a sequential prefetching technique that utilizes spatial locality and is straightforward to execute. This method triggers a prefetch for line $i+1$ upon accessing line i and encountering a cache miss. However, if line $i+1$ is already present in the cache, no memory access is initiated.